

WIP: Transparent Standby for Low-Power, Resource-Constrained Embedded Systems

A Programming Language-Based Approach

Anonymous Author(s)

Abstract

Standby efficiency for connected devices is one of the priorities of the G20's *Energy Efficiency Action Plan*. We propose transparent programming language mechanisms to enforce that applications remain in deepest standby modes for longest periods of time. We extend the synchronous programming language C  U with support for interrupt service routines and with a simple power management runtime. We developed device drivers based on these primitives on top of which applications can be built to take advantage of standby automatically. We also show that programs in C  U can keep a sequential structure to lower the barrier of adoption, even when applications require non-trivial concurrent behaviors.

CCS Concepts • **Computer systems organization** → **Embedded systems**; *Redundancy*; Robotics; • **Networks** → Network reliability;

Keywords ACM proceedings, L      , text tagging

ACM Reference Format:

Anonymous Author(s). 2018. WIP: Transparent Standby for Low-Power, Resource-Constrained Embedded Systems: A Programming Language-Based Approach. In *Proceedings of ACM SIGPLAN/SIGBED (LCTES'18)*. ACM, New York, NY, USA, 4 pages. https://doi.org/10.475/123_4

According to the International Energy Agency (IEA), the number of network-connected devices is expected to reach 50 billion by 2020 with the expansion of the Internet of Things (IoT) [6]. However, most of the energy to power these devices will be consumed in *standby mode*, i.e., when they are neither transmitting or processing data. For instance, standby power currently accounts for 10–15% of residential electricity consumption, and CO₂ emissions related to standby are equivalent to those of 1 million cars [6, 7]. The projected growth of IoT devices, together with the surprising effects of

standby consumption, made network standby efficiency one of the six pillars of the G20's *Energy Efficiency Action Plan*¹.

Given the projected scale of the IoT and the role of low-power standby towards energy efficiency, this paper has the following goals:

1. Address energy efficiency through extensive use of standby.
2. Target low-power, resource-constrained embedded architectures that form the IoT.
3. Provide standby mechanisms at the programming language level that scale to all applications.
4. Support transparent/non-intrusive standby mechanisms that reduce barriers of adoption.

Our approach lies at the bottom of the software development layers—programming language mechanisms—meaning that *all* applications take advantage of low-power standby modes automatically, without extra programming efforts. We extend the synchronous programming language C  U [9, 10] with support for interrupt service routines (ISRs) and with a simple power management runtime (PMR).

Each supported microcontroller requires bindings in C for the ISRs and PMR, and each peripheral requires a driver in C  U. These are write-once code that are typically packaged and distributed in a software development kit (SDK). Then, all new applications built on top of these drivers take advantage of standby automatically. As a proof of concept, we provide an open source SDK with support for 8-bit AVR/ATmega and 32-bit ARM/Cortex-M0 microcontrollers, and a variety of peripherals, such as GPIO, A/D converter, USART, SPI, and the nRF24L01 transceiver. We developed a number of simple applications using these peripherals concurrently and could verify that the applications remain in deepest standby modes for longest periods of time.

In Section 1, we compare the structure of programs in C  U and Arduino [2], whose primary goal is to reduce the barrier of adoption for non-technical users (e.g., designers and artists). We show that we can keep the intended sequential reasoning of Arduino even when applications require non-trivial concurrent behaviors. In Section 2, we discuss the software infrastructure that allows for unmodified programs in C  U to take advantage of standby automatically. In Section 3, we discuss future work and conclude the paper.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

LCTES'18, June 2018, Philadelphia, Pennsylvania, USA

   2018 Copyright held by the owner/author(s).

ACM ISBN 123-4567-24-567/08/06.

https://doi.org/10.475/123_4

¹G20's Energy Efficiency Action Plan: <https://www.iea-4e.org/projects/g20>

```

111 while (1) {
112     delay(1000);
113     int v =
114         analogRead();
115     radioWrite(v);
116 }

```

[a] Version in Arduino

```

1 loop do
2   await 1s;
3   var int v =
4     await AnalogRead();
5   await RadioWrite(v);
6 end

```

[b] Version in CéU

Figure 1. Sequence of I/O operations running in a loop.

1 The Structured Synchronous Programming Language CéU

CéU is a Esterel-based[9] reactive programming language targeting resource-constrained embedded systems [10]. It is grounded on the synchronous concurrency model, which has been successfully adopted in the context of hard real-time systems such as avionics and automobiles industry since the 80's [3]. The synchronous model trades power for reliability and has a simpler model of time that suits most requirements of IoT applications. On the one hand, this model cannot directly express time-consuming computations, such as compression and cryptography algorithms, which are typically either absent or delegated to auxiliary chips in the context of the IoT. On the other hand, all reactions to the external world are guaranteed to be computed in bounded time, ensuring that applications always reach an idle state amenable to standby mode. Overall, CéU aims to offer a concurrent, safe, and expressive alternative to C with the characteristics that follow:

Reactive: code only executes in reactions to events.

Structured: programs use structured control mechanisms, such as `await` (to suspend a line of execution), and `par` (to combine multiple lines of execution).

Synchronous: reactions run atomically and to completion on each line of execution, i.e., there's no implicit preemption or real parallelism.

Structured reactive programming lets developers write code in direct style, recovering from the inversion of control imposed by event-driven execution [1, 5, 8].

A Motivating Example

Figure 1.a shows a simple, easy-to-read program chunk in Arduino that executes forever in a loop a sequence of operations as follows: waits for 1 second (ln. 2), performs an A/D conversion (ln. 3–4), and broadcasts the value read (ln. 5). Figure 1.b shows the same chunk in CéU, with a noteworthy difference that operations that interact with the environment and take time use the `await` keyword. The traditional structured paradigm encouraged in Arduino (with blocks, loops, and sequences) allows for simple and readable code, avoiding the complexity of dealing with ISRs. However, the

```

uint32_t prv = 1 par/or do
  millis(); 2   await RadioAval();
while (1) { 3   with
  if (radioAval()) { 4   loop do
    break; 5     await 1s;
  } 6     var int v =
  uint32_t cur = 7       await AnalogRead();
  millis(); 8       await RadioWrite(v);
  if (cur>prv+1000) { 9   end
    prv = cur; 10  end
    int v = 11
      analogRead(); 12
      radioWrite(v); 13
    } 14
  } 15
}

```

[a] Version in Arduino

[b] Version in CéU

Figure 2. Achieving concurrency between I/O operations.

use of blocking operations, such as `delay(1000)`, prevents that other operations execute concurrently.

Suppose we now want that, at any time, receiving a message via radio should immediately abort the loop in Figure 1.a. Since the message might arrive concurrently with any of the blocking operations, we need to change the structure of the program. Figure 2.a changes the blocking operation `delay` to the polling operation `millis`, which immediately returns the number of milliseconds since the reset. Now, we start by registering the current time (ln. 1–2) and, on each loop iteration, we recheck the time to see if one second has elapsed (ln. 7–9). Since these operations are non-blocking, we can intercalate the execution with checks for message arrivals (ln. 4–6). If the time is up, we start counting it again (ln. 10) before proceeding to the original operations in sequence (ln. 11–13). The original structured style has been drastically violated to accommodate concurrency. In the example, we only adapted the `delay` operation, but the other blocking operations (`analogRead` and `radioWrite`) would also need to be changed to achieve maximum concurrency. Alternatively, we could resort to ISRs or implement an event-driven scheduler to handle the operations [4], but ultimately, the program readability would still be compromised.

The program in Figure 2.b in CéU extends the one in Figure 1.b to accommodate concurrency. The original code remains unmodified (ln. 4–9) and concurrency is achieved through the `par/or` construct, which creates two lines of execution and terminates when either of them terminates, aborting the other automatically. This approach preserves the sequential, easy-to-read style while accommodating concurrency seamlessly.

Standby Considerations

The structure of the program in Figure 2.b also indicates which peripherals are active at a given time. For instance, when the program is awaiting concurrently in lines 2 and 7, only the radio transceiver and A/D converter can awake the program. Hence, the language runtime can choose the most energy-efficient sleep mode that allows these peripherals to awake the microcontroller from associated interrupts. Since the semantics of C    enforces the program to always reach `await` statements in all active lines of execution, it is always possible to put the microcontroller into the optimal sleep mode after each reaction to the environment.

2 Standby Infrastructure

In order to empower the example in Figure 2.b with automatic standby, we had to make some modifications and extensions to C   :

- Modify the runtime of C    to be interrupt driven and to put the microcontroller in standby after each reaction to the environment.
- Provide operations for the drivers to indicate which interrupts can awake the program.
- Support ISRs in C    to generate input events to the program.

Figure 3 shows the driver for the A/D converter in C   . This code is specific to the *ATmega328p* microcontroller and must be adapted to work in other platforms. For simplicity, we assume in the paper that the converter has a single channel to avoid having to deal with multiplexing.

The driver exposes raw I/O events (ln. 3–4) that only deal with low-level port manipulation in the microcontroller. Output events are triggered with the `emit` keyword (ln. 29), while input events are captured with the `await` keyword (ln. 30). The output event `ADC_REQUEST` (ln. 9–15) enables ADC interrupts and starts an analog-to-digital conversion in the single channel `A0`. Any code in between `{` and `}` is treated as an inline C chunk, allowing for easy integration with C for low-level operations.

The `async/isr` construct of C    defines an ISR which executes asynchronously with the program when the specified interrupt occurs. Only ISRs can emit input events to the program. In the example, we define an ISR to handle ADC interrupts which fire whenever a conversion is complete (ln. 17–21). Although the ISR block executes asynchronously on interrupts, the input emission (ln. 20) only takes effect on a subsequent reaction, when the synchronous part of the program becomes idle. This way, race conditions are only possible within `async/isr` blocks, which are typically hidden inside device drivers. C    also provides an atomic primitive to protect critical sections of code.

The low-level events are the pieces that vary among platforms. A driver can also offer a higher-level portable abstraction to be exposed to client code. In the example, the

```

1 // Exposed driver functionality
2
3 output void ADC_REQUEST; // low-level request
4 input int ADC_DONE; // low-level response
5 code AnalogRead (void) -> int; // high-level abstraction
6
7 // Driver implementation
8
9 output void ADC_REQUEST do
10 {
11     ADMUX = 0x40 | (A0 & 0x07);
12     bitSet(ADCSRA, ADIE); // enables interrupt
13     bitSet(ADCSRA, ADSC); // starts the conversion
14 }
15 end
16
17 async/isr {ADC_vect_num} do
18 { bitClear(ADCSRA, ADIE); } // disables interrupt
19 var int value = {ADC}; // reads register with the value
20 emit ADC_DONE(value);
21 end
22
23 code AnalogRead (void) -> int do
24 {PM_SET(PM_ADC, 1);}
25 do finalize with
26 {PM_SET(PM_ADC, 0);}
27 end
28
29 emit ADC_REQUEST(pin);
30 var int value = await ADC_DONE;
31
32 escape value;
33 end

```

Figure 3. C    driver for the *ATmega328p* A/D converter.

`AnalogRead` abstraction (ln. 23–33) takes care of starting and awaiting the conversion (ln. 29–30), as well as dealing with the power management runtime (PMR). The `PM_SET(PM_ADC, 1)` (ln. 24) tells the system that, when entering in sleep mode, the ADC must be kept running. The `PM_SET(PM_ADC, 0)` inside the `finalize` clause (ln. 25–27) releases the ADC subsystem.

The `finalize` construct of C    executes the nested code whenever the enclosing block terminates or is aborted externally. The example of Figure 2.b invokes the `AnalogRead` abstraction (ln. 7) concurrently with `RadioAvail` (ln. 2). It is possible that the `AnalogRead` terminates normally or that a radio message arrives during the A/D conversion, which causes the `AnalogRead` to abort abruptly. In either case, the `finalize` clause executes and puts the PMR in a consistent state.

The PMR also expects a platform-specific power management module to be able to put the microcontroller into

```

331 1 #define PM_GET(dev)    bitRead(pm,dev)
332 2 #define PM_SET(dev,v) bitWrite(pm,dev,v)
333 3
334 4 static u32 pm = 0;
335 5
336 6 enum {
337 7     CEU_PM_ADC = 0,
338 8     CEU_PM_TIMER1,
339 9     <...>,
340 10 };
341 11
342 12 void pm_sleep (void) {
343 13     if (PM_GET(PM_TIMER1) || <...>) {
344 14         LowPower.idle(PM_GET(PM_ADC),<...>)
345 15     } else if (PM_GET(PM_ADC)) {
346 16         LowPower.adcNoiseReduction(<...>);
347 17     } else {
348 18         LowPower.powerDown(<...>);
349 19     }
350 20 }
351 21 }

```

Figure 4. Power management module for the *ATmega328p* microcontroller.

the most efficient sleep mode possible. Figure 4 implements the `pm_sleep` function for the *ATmega328p* microcontroller which the PMR calls when the program is idle. Each device has an associated index (ln. 6–10) in the `pm` bit vector (ln. 4). The driver manipulates its device’s index to indicate its state (Figure 3, ln. 24,26). The `pm_sleep` queries the vector to choose the appropriate sleep mode. In the example, if the timer is active (ln. 13), the microcontroller can only use the least efficient mode² (ln. 14). In the best case, e.g., if only external interrupts are required, the microcontroller can use the most efficient mode (ln. 18).

With all the standby infrastructure set, the program in Figure 2.b will automatically take advantage of the longest periods of time and deepest sleep modes possible.

3 Conclusion and Future Work

In this work, we address standby efficiency for embedded devices at the level of programming languages. We propose a software infrastructure for the programming language C  u that encompass a power management runtime and support for interrupt service routines in the language. Our approach relies on the synchronous semantics of the language which enforces that reactions to the environment always reach an idle state amenable to standby. We then take energy-unaware programs in C  u and show that, using the proposed

infrastructure, they can take advantage of the longest periods of time and deepest sleep modes possible.

In future work, in order to evaluate the gains in energy efficiency with the proposed infrastructure, we will evaluate the consumption of realistic applications. The Arduino community has an abundance of open-source projects which can be rewritten in C  u take advantage of transparent standby modes. In this scenario, we can evaluate the time to rewrite, the resulting program structure, and the actual energy consumption efficiency.

References

- [1] A. Adya et al. 2002. Cooperative Task Management Without Manual Stack Management. In *Proceedings of ATEC’02*. USENIX Association, 289–302.
- [2] Massimo Banzi and Michael Shiloh. 2014. *Getting started with Arduino: the open source electronics prototyping platform*. Maker Media, Inc.
- [3] Albert Benveniste, Paul Caspi, Stephen A. Edwards, Nicolas Halbwachs, Paul Le Guernic, and Robert De Simone. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.
- [4] David Gay, Philip Levis, Robert von Behren, Matt Welsh, Eric Brewer, and David Culler. 2003. The nesC Language: A Holistic Approach to Networked Embedded Systems. In *Proceedings of PLDI’03*. 1–11.
- [5] Ingo Maier, Tiark Rompf, and Martin Odersky. 2010. *Deprecating the observer pattern*. Technical Report.
- [6] OECD/IEA. 2014. *More Data Less Energy—Making Network Standby More Efficient in Billions of Connected Devices*. Technical Report. International Energy Agency.
- [7] Australian Greenhouse Office. 2002. *Money isn’t all you’re saving: Australia’s standby power strategy 2002–2012*. Australian Greenhouse Office. 23 pages.
- [8] Guido Salvaneschi et al. 2014. REScala: Bridging between object-oriented and functional style in reactive applications. In *Proceedings of Modularity’13*. ACM, 25–36.
- [9] Francisco Sant’anna, Roberto Ierusalimsky, Noemi Rodriguez, Silvana Rossetto, and Adriano Branco. 2017. The Design and Implementation of the Synchronous Language C  u. *ACM Trans. Embed. Comput. Syst.* 16, 4, Article 98 (July 2017), 26 pages. <https://doi.org/10.1145/3035544>
- [10] Francisco Sant’Anna, Noemi Rodriguez, Roberto Ierusalimsky, Olaf Landsiedel, and Philippas Tsigas. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys’13*. ACM.

² We use an external library for the sleep modes: <http://www.rocketstream.com/blog/2011/07/04/lightweight-low-power-arduino-library/>